
Conclusion

— **Colin Lockard**, Prashant Shiralkar, —
Xin Luna Dong, Hannaneh Hajishirzi



PAUL G. ALLEN SCHOOL
OF COMPUTER SCIENCE & ENGINEERING

Four Challenges

1. Diversity of data
2. Multiple modalities of text
3. Lack of training data
4. Unknown unknowns

Can we build a single extractor to find **consistent signals** across these diverse elements of data **across all modalities of text**?

Key Intuitions

- Diversity: Identifying consistent patterns
 - Leverage consistency in model/representation
 - Combining information from multiple modalities can give more consistent signals
- Lack of training data: Learning with limited labels
 - Find automated ways to label data
 - Employ weak or semi-supervision in limited labeled data settings
- Unknown unknowns: Stay open--Sacrificing granularity of knowledge representation allows for easier scaling

Unstructured Text: Short Answers

- **Consistency**
 - Model problem as text span classification and relationships between spans
 - Word embedding models help capture text semantics
- **Training data**
 - Weak supervision gives cheap training data
- **OpenIE**
 - Discovery of new types and relationships

Semi-Structured Text: Short Answers

- **Consistency**
 - Leverage general key-value pair consistency universal in templates
 - Leverage site-level consistency in layout and presentation
- **Training data**
 - Use distant supervision to generate cheap, but noisy training data
- **OpenIE**
 - Discover new relations by label propagation

Tabular text - Short Answers

- **Subject column detection**
 - Leverage generic features of subject entities such as value uniqueness, string type, number of characters and words
- **Column class detection**
 - Leverage external data -- web extracted triples, knowledge graph
- **Relation extraction between column pair**
 - Measure similarity between a column and entities of a type in a knowledge base

Multi-modal extraction: Short answers

- **Diversity**
 - Textual, layout, and visual signals can combine to form consistent patterns
- **Training data**
 - Multi-modal signals allow for accurate and easy creation of training data with Data Programming
- **OpenIE**
 - Visual semantics help make OpenIE extractions from semi-structured documents without prior knowledge of the subject domain

Future Directions - Unstructured text

- Full document understanding
 - Relation extraction beyond single sentence/paragraph
- Faster embedding models for scalability
- Non-English languages

Future Directions - Semi-structured text

- N-ary relations
- Relations not involving page topic

Future Directions - Tabular text

- Direct extraction (not relying on existing knowledge)

Future Directions - Multi-modal extraction

- Combine all signals from a document
- Make use of images
- Operate from jpgs, scanned pdfs
- Pre-training webpage representations
- Automated ontology construction
- Reproducible research
 - Webpage visual features depend on browser, CSS/JS availability, etc.

Outline

- Introduction (40 minutes)
- Part 1a: Unstructured Text (25 minutes)
- Part 1b: Unstructured Text: Methods (10 minutes)
- Live Q&A (15 minutes)
- Break (30 minutes)
- Part 2: Semi-structured and Tabular Text (40 minutes)
- Part 3: Multi-modal Extraction and Conclusion (35 minutes)
- **Live Q&A (15 minutes)**

Thank you!

[https://sites.google.com/view/
acl-2020-multi-modal-ie](https://sites.google.com/view/acl-2020-multi-modal-ie)

Please join us in
the Zoom Chat!
